

# Samash technical Test : Student Data Analytics

In [17]:

```
import pandas as pd
import numpy as np
import csv
from datetime import datetime
import seaborn as sns
import matplotlib.pyplot as plt
```

In [18]:

```
dataSet = pd.read_csv("student_data.csv")
```

In [19]:

```
dataSet.head()
```

Out[19]:

|   | Student ID | Student Name    | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization   | Fees   | Discount on Fees |
|---|------------|-----------------|---------------|------------------------|-------------------|-----------------------------|------------------|------------------|--------|------------------|
| 0 | 165527     | Bryan Rogers    | 2006-01-19    | Computer Science       | 2020              | 2017                        | 3                | Web Development  | 155152 | 19572            |
| 1 | 635763     | James Hogan     | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning | 157870 | 14760            |
| 2 | 740021     | David Robinson  | 1997-12-02    | Civil Engineering      | 2017              | 2022                        | 1                | Network Security | 55662  | 5871             |
| 3 | 433076     | Susan Miller    | 1999-10-30    | Computer Science       | 2021              | 2019                        | 1                | Data Science     | 134955 | 17284            |
| 4 | 441628     | Brittany Martin | 1998-01-10    | Chemical Engineering   | 2016              | 2018                        | 1                | Network Security | 125934 | 14871            |

In [20]:

```
dataSet.shape
```

Out[20]:

```
(200000, 10)
```

In [21]:

```
dataSet.columns
```

Out[21]:

```
Index(['Student ID', 'Student Name', 'Date of Birth', 'Field of Study',
      'Year of Admission', 'Expected Year of Graduation', 'Current Semester',
      'Specialization', 'Fees', 'Discount on Fees'],
      dtype='object')
```

In [22]:

```
dataSet.dtypes
```

Out[22]:

```
Student ID          int64
Student Name        object
Date of Birth        object
Field of Study       object
Year of Admission    int64
Expected Year of Graduation  int64
Current Semester     int64
Specialization       object
Fees                 int64
Discount on Fees     int64
dtype: object
```

In [23]:

```
missing_values = dataSet.isnull().sum()
missing_values
```

Out[23]:

```
Student ID          0
Student Name        0
Date of Birth        0
Field of Study       0
Year of Admission    0
Expected Year of Graduation  0
Current Semester     0
Specialization       0
Fees                 0
Discount on Fees     0
dtype: int64
```

## Cleaning & data quality

## Parsing dates

In [26]:

```
#Verify the date formate using : to_date and handelng errors
#The errors = 'coerce' parameter handles any invalid or unparsable dates by converting them to NaT (Not a Time)
#We would Like to have this formate %Y-%m-%d for all the data attributes

def clean_dates(data, col):
    data[col] = pd.to_datetime(data[col], errors='coerce')
    return data

cleaned_data = clean_dates(dataSet, 'Date of Birth')

cleaned_data.head()
cleaned_data.dtypes
```

Out[26]:

```
Student ID          int64
Student Name        object
Date of Birth       datetime64[ns]
Field of Study      object
Year of Admission   int64
Expected Year of Graduation int64
Current Semester    int64
Specialization      object
Fees                int64
Discount on Fees    int64
dtype: object
```

In [27]:

```
def clean_parse_string(data):
    object_columns = data.select_dtypes(include='object').columns
    data[object_columns] = data[object_columns].apply(lambda x: x.astype(str))
    return data

cleaned_data = clean_parse_string(cleaned_data)
cleaned_data.infer_objects().dtypes
```

Out[27]:

```
Student ID          int64
Student Name        object
Date of Birth       datetime64[ns]
Field of Study      object
Year of Admission   int64
Expected Year of Graduation int64
Current Semester    int64
Specialization      object
Fees                int64
Discount on Fees    int64
dtype: object
```

## Adjusting admission year and graduation year column values

In [28]:

```
# Function to compare Admission and graduation Years. the second column (graduation Year) must contain greater values
def compare_columns(df, column1, column2, result_column):
    df[result_column] = df[column1] > df[column2]
    mask = df[column1] > df[column2]
    df.loc[mask, [column1, column2]] = df.loc[mask, [column2, column1]].values
    return df

# Compare columns and add the result to a new column
compared_data = compare_columns(dataSet, 'Year of Admission', 'Expected Year of Graduation', 'result_column')

# Print the modified DataFrame
compared_data.head(2)
```

Out[28]:

|   | Student ID | Student Name | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization   | Fees   | Discount on Fees | result |
|---|------------|--------------|---------------|------------------------|-------------------|-----------------------------|------------------|------------------|--------|------------------|--------|
| 0 | 165527     | Bryan Rogers | 2006-01-19    | Computer Science       | 2017              | 2020                        | 3                | Web Development  | 155152 | 19572            |        |
| 1 | 635763     | James Hogan  | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning | 157870 | 14760            |        |

## Getting Expected Education duration for each student

In [29]:

```
df=compared_data
df['Expected Student Education Duration'] = df['Expected Year of Graduation'] - df['Year of Admission'] + 1
df.head(2)
```

Out[29]:

|   | Student ID | Student Name | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization   | Fees   | Discount on Fees | result |
|---|------------|--------------|---------------|------------------------|-------------------|-----------------------------|------------------|------------------|--------|------------------|--------|
| 0 | 165527     | Bryan Rogers | 2006-01-19    | Computer Science       | 2017              | 2020                        | 3                | Web Development  | 155152 | 19572            |        |
| 1 | 635763     | James Hogan  | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning | 157870 | 14760            |        |

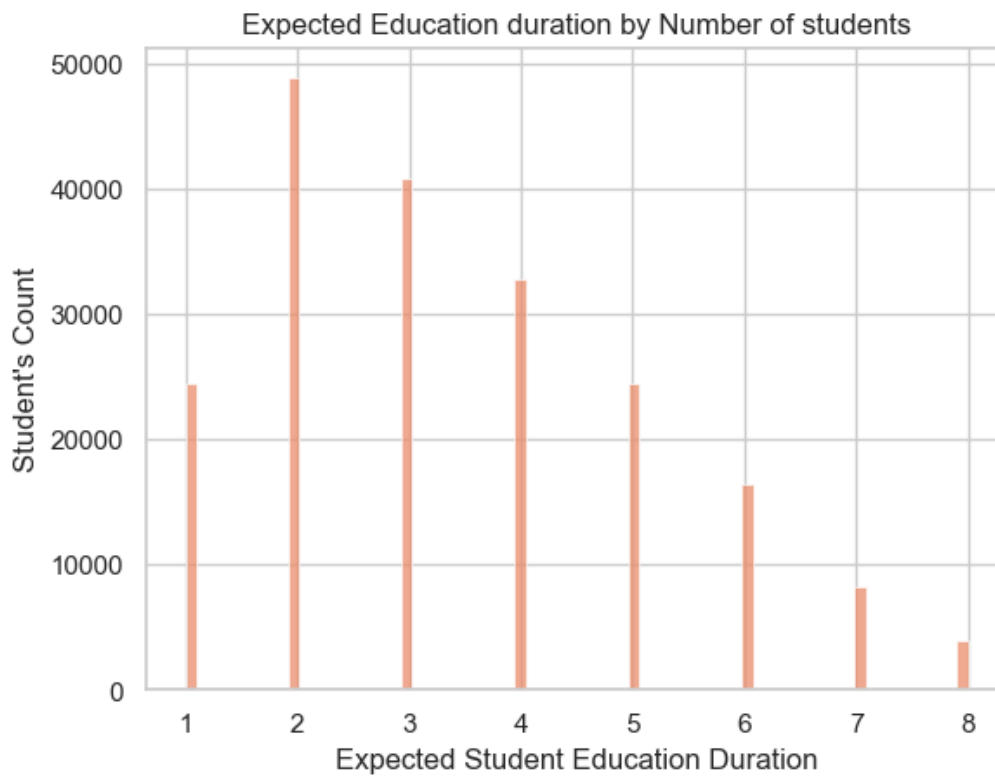
In [74]:

```
import matplotlib.pyplot as plt
sns.set_theme(style="whitegrid", palette="flare")
sns.histplot(data=df, x="Expected Student Education Duration")
plt.title("Expected Education duration by Number of students")
axes = plt.gca()
axes.set_ylabel("Student's Count")

"""for bar in axes.patches:
    # Get the height of the bar
    height = bar.get_height()

    # Add text annotation on top of the bar
    axes.annotate(f'{height:.2f}',
                  (bar.get_x() + bar.get_width() / 2, height),
                  ha='center', va='bottom')"""

plt.show()
```



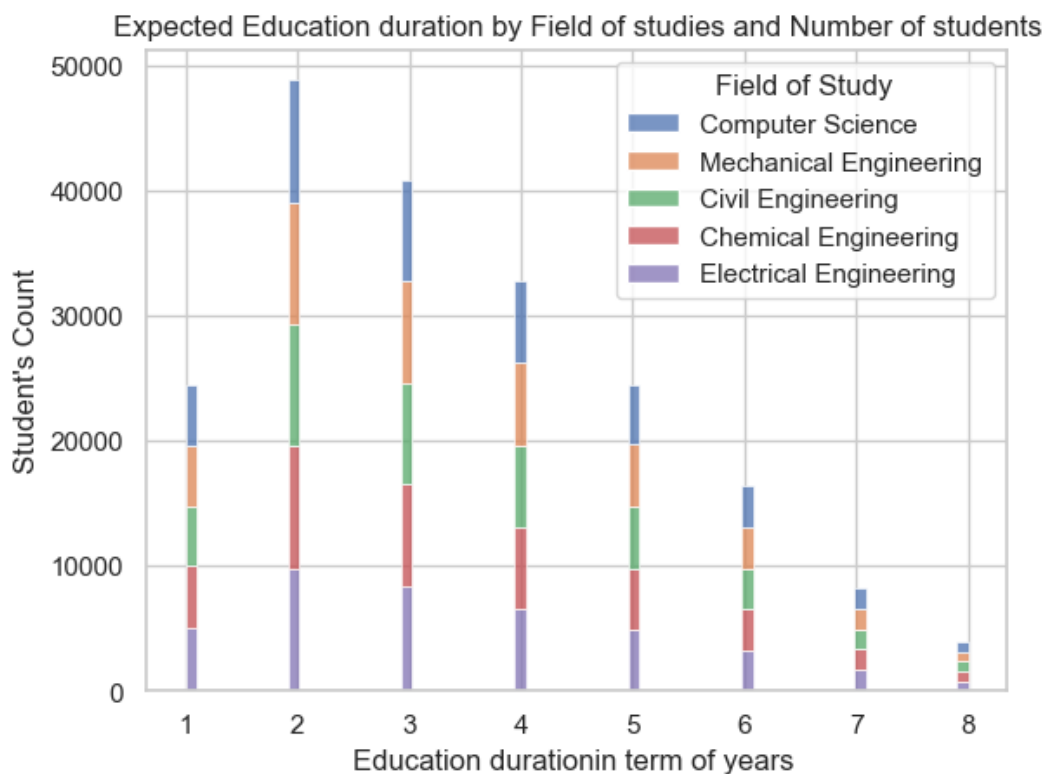
In [62]:

```
import matplotlib.pyplot as plt
sns.set_theme(style="whitegrid", palette="deep")
sns.histplot(data=df, x="Expected Student Education Duration", hue="Field of Study", multiple="stack")
axes = plt.gca()
axes.set_ylabel("Student's Count")
axes.set_xlabel("Education durationin term of years ")
plt.title("Expected Education duration by Field of studies and Number of students")

plt.show
```

Out[62]:

```
<function matplotlib.pyplot.show(close=None, block=None)>
```



## Adjusting semester dispatching for student expected to graduate the same admission year

In [198]:

```
#We noticed that in some rows, the year of admission equals year of graduation
filtered_df = df[df['Year of Admission'] == df['Expected Year of Graduation']]
filtered_df.head(2)
```

Out[198]:

|   | Student ID | Student Name | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization          | Fees   | Discount on Fees | result |
|---|------------|--------------|---------------|------------------------|-------------------|-----------------------------|------------------|-------------------------|--------|------------------|--------|
| 1 | 635763     | James Hogan  | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning        | 157870 | 14760            |        |
| 6 | 268329     | Erica Owens  | 2003-06-20    | Mechanical Engineering | 2020              | 2020                        | 1                | Artificial Intelligence | 52994  | 5231             |        |

In [199]:

```
#we can notice that if in the same year the student's current semester is greater then 2, the value is 1  
#to minimize the error, if the expected year isn't 2023 then we set the value at 2 (graduated)
```

```
df.loc[df['Year of Admission'] == df['Expected Year of Graduation'] , 'Current Semester'] = 2  
df.head(2)
```

Out[199]:

|   | Student ID | Student Name | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization   | Fees   | Discount on Fees | result |
|---|------------|--------------|---------------|------------------------|-------------------|-----------------------------|------------------|------------------|--------|------------------|--------|
| 0 | 165527     | Bryan Rogers | 2006-01-19    | Computer Science       | 2017              | 2020                        | 3                | Web Development  | 155152 | 19572            |        |
| 1 | 635763     | James Hogan  | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning | 157870 | 14760            |        |

In [132]:

```
#Since we're in the first semester of 2023, we turn the value of undergraduated student to semester 1  
duplication = filtered_df [filtered_df ['Expected Year of Graduation'] == 2023]  
print(duplication)  
#there's no such condition, we won't apply changes on df which is the copy of the data source
```

Empty DataFrame

Columns: [Student ID, Student Name, Date of Birth, Field of Study, Year of Admission, Expected Year of Graduation, Current Semester, Specialization, Fees, Discount on Fees, result\_column, Expected Student Education Duration]

Index: []

In [133]:

```
#the portion of student expected to graduate the same year they are admitted in the total data set
```

```
portion = len(filtered_df) / len(df)  
print("portion shape is", filtered_df.shape)  
portionPourcentage = len(filtered_df) / len(df) * 100  
print("pourcentage of duplication", portionPourcentage)
```

portion shape is (24498, 12)

pourcentage of duplication 12.249

In [81]:

```
#We will assume that the proportion of students within each field of study  
#who are expected to graduate in the same year they were admitted  
  
grouped_df = df.groupby('Field of Study').apply(lambda x: len(x[x['Year of Admission']  
                                                    == x['Expected Year of Graduation']]) / len(x))  
print(" Pourcentage of One Year duration field compared to the other periods per field: \n", grouped_df)  
grouped_df.plot.pie(autopct='%1.1f%%', startangle=90)  
plt.title('One Year duration programs by field')
```

Pourcentage of One Year duration field compared to the other periods per field:


| Field of Study         |           |
|------------------------|-----------|
| Chemical Engineering   | 12.241379 |
| Civil Engineering      | 12.093327 |
| Computer Science       | 12.389890 |
| Electrical Engineering | 12.531944 |
| Mechanical Engineering | 11.985178 |

dtype: float64

Out[81]:

Text(0.5, 1.0, 'One Year duration programs by field')

One Year duration programs by field





## Pourcentage of ungraduate students (still in Uni) and their age

In [82]:

```
# Selecting rows with graduation year >= 2023
age = df[df['Expected Year of Graduation'] >= 2023]
# get current year
current_year = datetime.now().year
age['Age of Undergraduated student'] = current_year - pd.DatetimeIndex(age['Date of Birth']).year
age.head(2)
#age.shape
```

C:\Users\asus\AppData\Local\Temp\ipykernel\_8836\1112744852.py:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
age['Age of Undergraduated student'] = current_year - pd.DatetimeIndex(age['Date of Birth']).year
```

Out[82]:

|    | Student ID | Student Name  | Date of Birth | Field of Study       | Year of Admission | Expected Year of Graduation | Current Semester | Specialization          | Fees   | Discount on Fees | rest |
|----|------------|---------------|---------------|----------------------|-------------------|-----------------------------|------------------|-------------------------|--------|------------------|------|
| 18 | 439048     | Gary Garcia   | 2005-07-21    | Chemical Engineering | 2020              | 2023                        | 3                | Machine Learning        | 88470  | 292              |      |
| 22 | 400958     | Nicole Grimes | 1998-11-30    | Civil Engineering    | 2021              | 2023                        | 4                | Artificial Intelligence | 152798 | 19868            |      |

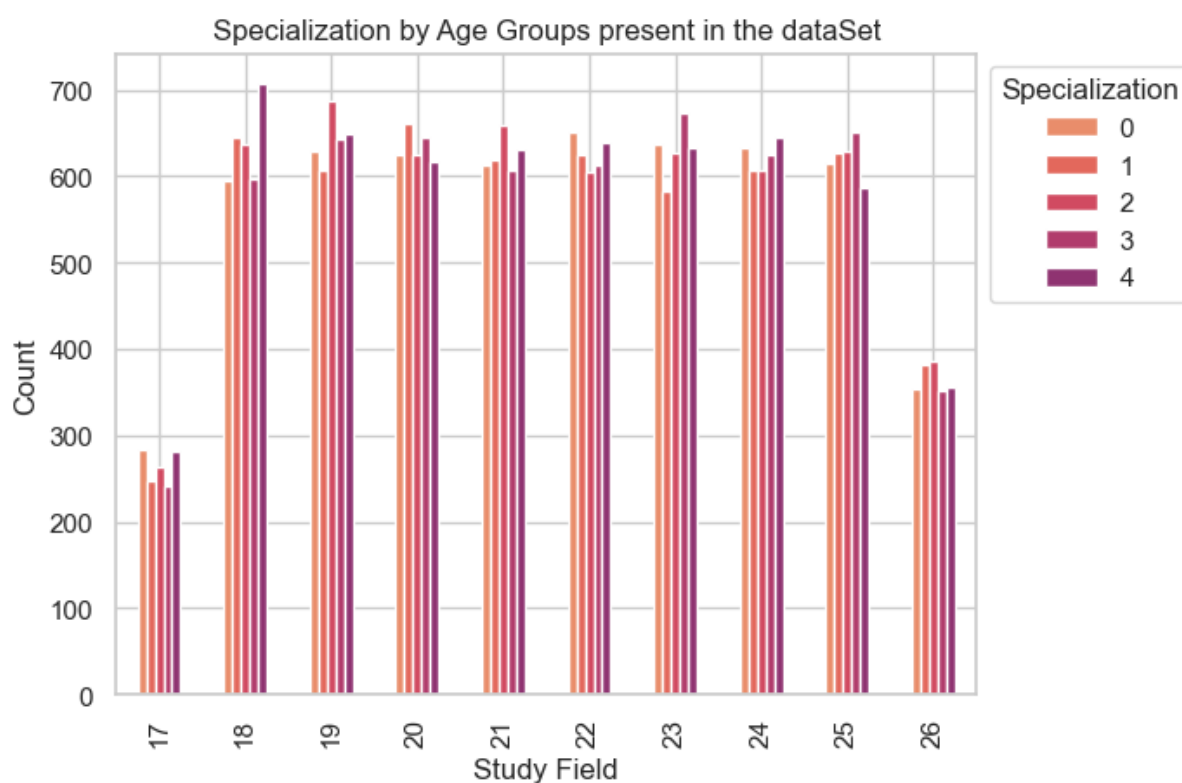
In [165]:

```
grouped_students = age.groupby('Age of Undergraduated student')[ "Specialization"].value_counts().unstack
ax = grouped_students.plot(kind='bar')

# Set the x-axis Label
plt.xlabel('Study Field')

# Set the y-axis Label
plt.ylabel('Count')

# Set the title of the plot
plt.title('Specialization by Age Groups present in the dataSet')
sns.move_legend(ax, "upper left", bbox_to_anchor=(1, 1))
# Display the plot
plt.show()
```



## Graduate vs Ungraduate students pourcentage

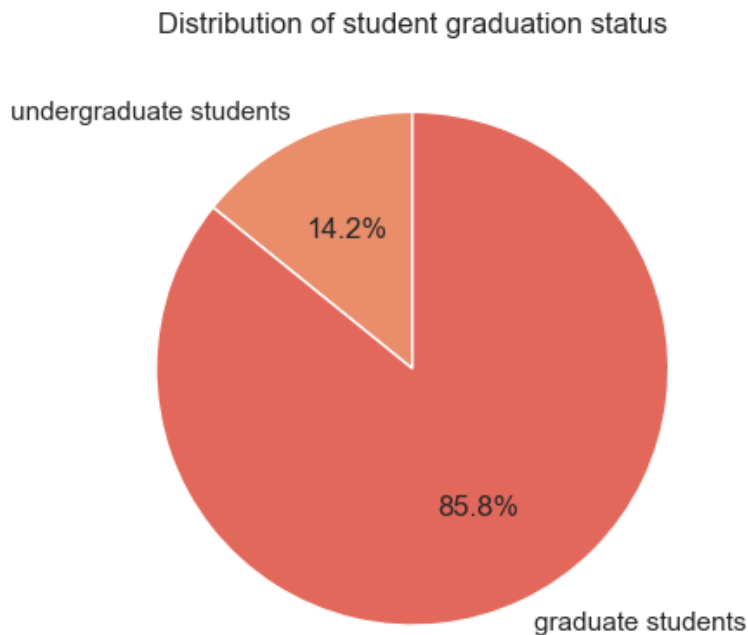
In [108]:

```
# we want to see the diffrent pourcentage of student graduation statu in the total data
undergraduate_Pourcentage = len(age) / len(df) * 100
graduate_Pourcentage = 100 - undergraduate_Pourcentage
print("pourcentage of undergraduate student present in the dataset", undergraduate_Pourcentage,"%")
print("pourcentage of graduated student present in the dataset",graduate_Pourcentage,"%")
```

pourcentage of undergraduate student present in the dataset 14.180000000000001 %  
pourcentage of graduated student present in the dataset 85.82 %

In [109]:

```
plt.pie([undergraduate_Pourcentage,graduate_Pourcentage], labels=['undergraduate students','graduate students'],
plt.title('Distribution of student graduation status')
plt.show()
```



## Fees analysis based on Field of studies using Kmeans

In [202]:

```
unique_fields= df['Field of Study'].unique()
print('Number of main fields of studies is : ', len(unique_fields))
unique_fields
```

Number of main fields of studies is : 5

Out[202]:

```
array(['Computer Science', 'Mechanical Engineering', 'Civil Engineering',
       'Chemical Engineering', 'Electrical Engineering'], dtype=object)
```

In [203]:

```
df['Fees after reduction'] = df['Fees'] - df['Discount on Fees']
df.head(2)
```

Out[203]:

|   | Student ID | Student Name | Date of Birth | Field of Study         | Year of Admission | Expected Year of Graduation | Current Semester | Specialization   | Fees   | Discount on Fees | result |
|---|------------|--------------|---------------|------------------------|-------------------|-----------------------------|------------------|------------------|--------|------------------|--------|
| 0 | 165527     | Bryan Rogers | 2006-01-19    | Computer Science       | 2017              | 2020                        | 3                | Web Development  | 155152 | 19572            |        |
| 1 | 635763     | James Hogan  | 1999-05-23    | Mechanical Engineering | 2020              | 2020                        | 2                | Machine Learning | 157870 | 14760            |        |

In [112]:

```
from sklearn.cluster import KMeans

from sklearn.preprocessing import LabelEncoder

#string to numeric value

def convert_string_columns(data):

    encoded_data = data.copy()

    for column in encoded_data.select_dtypes(include='object'):

        label_encoder = LabelEncoder()

        encoded_data[column] = label_encoder.fit_transform(encoded_data[column])

    return encoded_data

# date to numeric value
try :
    age['Reduction pourcentage'] = age['Discount on Fees'] / age['Fees'] * 100
except:
    print("problem of assignment")

age['Date of Birth'] = pd.to_datetime(df['Date of Birth'])

age['Date of Birth'] = (pd.to_datetime('2023-06-17') - age['Date of Birth'] ).dt.total_seconds().abs()

age = convert_string_columns(age)

age.head(2)
```

Out[112]:

|    | Student ID | Student Name | Date of Birth | Field of Study | Year of Admission | Expected Year of Graduation | Current Semester | Specialization | Fees   | Discount on Fees | re |
|----|------------|--------------|---------------|----------------|-------------------|-----------------------------|------------------|----------------|--------|------------------|----|
| 18 | 439048     | 8435         | 565056000.0   | 0              | 2020              | 2023                        | 3                | 2              | 88470  | 292              |    |
| 22 | 400958     | 18189        | 774576000.0   | 1              | 2021              | 2023                        | 4                | 0              | 152798 | 19868            |    |

In [163]:

```
kmeans = KMeans(n_clusters=5)

kmeans.fit(age)

kmeans.fit(age[['Age of Undergraduated student', 'Reduction pourcentage', 'Field of Study']])

# Get the cluster Labels
cluster_labels = kmeans.labels_

# Add the cluster Labels to the data frame
age['Cluster'] = cluster_labels

age.head()
```

C:\ProgramData\anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning  
warnings.warn(  
C:\ProgramData\anaconda3\lib\site-packages\sklearn\cluster\\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning  
warnings.warn(  
Out[163]:

|    | Student ID | Student Name | Date of Birth | Field of Study | Year of Admission | Expected Year of Graduation | Current Semester | Specialization | Fees   | Discount on Fees | res |
|----|------------|--------------|---------------|----------------|-------------------|-----------------------------|------------------|----------------|--------|------------------|-----|
| 18 | 439048     | 8435         | 565056000.0   | 0              | 2020              | 2023                        | 3                | 2              | 88470  | 292              |     |
| 22 | 400958     | 18189        | 774576000.0   | 1              | 2021              | 2023                        | 4                | 0              | 152798 | 19868            |     |
| 25 | 328657     | 18036        | 767836800.0   | 4              | 2019              | 2023                        | 1                | 2              | 87936  | 2612             |     |
| 28 | 989259     | 19293        | 813196800.0   | 1              | 2019              | 2023                        | 2                | 2              | 135043 | 993              |     |
| 38 | 200648     | 7856         | 562032000.0   | 4              | 2018              | 2023                        | 2                | 3              | 194644 | 36082            |     |

In [128]:

```
df_total = dataSet[['Student ID', 'Field of Study']]
df_age = age[['Student ID', 'Field of Study']]

# Merge the selected data frames based on ID
result_merge = pd.merge(df_total, df_age, on='Student ID', how='inner')
sorted_df = result_merge.sort_values('Field of Study_x')
sorted_df = sorted_df.drop_duplicates(subset='Field of Study_x')
sorted_df
```

Out[128]:

|       | Student ID | Field of Study_x       | Field of Study_y |
|-------|------------|------------------------|------------------|
| 0     | 439048     | Chemical Engineering   | 0                |
| 19737 | 957902     | Civil Engineering      | 1                |
| 26593 | 991070     | Computer Science       | 2                |
| 9190  | 527903     | Electrical Engineering | 3                |
| 13019 | 849558     | Mechanical Engineering | 4                |

In [129]:

```
sns.scatterplot(data=age, x='Age of Undergraduated student', y='Reduction pourcentage',
               size = "Cluster", sizes=(20, 200), hue='Cluster', palette='deep')

min_age = age['Age of Undergraduated student'].min()

max_age = age['Age of Undergraduated student'].max()

plt.xlim(xmin=min_age - 1, xmax=max_age + 1, emit=True, auto=False)

plt.show()
```

